

# CMPT 300

## Assignment 2

### Our new shell - cshell

Marks: 80 marks (plus 20 bonus marks)

#### Notes:

1. Failure to follow the instructions and rules may lead to failed test cases and/or a final grade of 0.
2. You can do this assignment individually or in a team of two. If you are doing it in a group, only one submission per group is required. For example:

If Person A and B are in one group, only one person has to accept the assignment link and do the submission.

Person A: upload all files (.txt, .c), group.txt.

Person B: no submission required.

For group assignment, submit a "**group.txt**" indicating the Github ids of each member. For example: if I am working with John in a group then group.txt will be as below:

```
hazra.imran  
Johndoe
```

3. You may submit multiple times until the deadline. Grade penalties will be imposed for late submissions (see the course outline for the details).
4. Always plan before coding.
5. All the codes in this lab must be done using C language only. No other languages should be used.
6. Use function-level and inline comments throughout your code. We will not be specifically grading documentation. However, remember that you will not be able to comment on your code unless sufficiently documented. Take the time to document your code as you develop it properly.
7. We will carefully analyze the code submitted to look for plagiarism signs, so please do not do it! If you are unsure about what is allowed, please talk to an instructor or a TA.

#### Coding Rules

- You have to follow the file name as specified in the instructions.
- **Makefile:** Makefile provides the following functionality:
  - **all:** compiles your program (this is the default behavior), producing an executable file named the same as the C file.
  - **clean:** deletes the executable file and any intermediate files (.o, specifically)
  - You will receive 0 if your makefile fails.
    - o Check your build to ensure that there are no errors.
    - o Visit TA's programming office hours to get help.

## Assignment Goals

- To understand the relationship between OS command interpreters (shells), system calls, and the kernel.
- To design and implement an extremely simple shell and system call (Bonus).

For this assignment, you should understand the concepts of environment variables, system calls, standard input and output, I/O redirection, parent and child processes, current directory, pipes, jobs, foreground and background, signals, and end-of-file.

## Background

OS command interpreter is the program that people interact with to launch and control programs. On UNIX systems, the command interpreter is usually called the **shell**. It is a user-level program that gives people a command-line interface to launch, suspend, or kill other programs. sh, ksh, csh, tcsh, bash, ... are all examples of UNIX shells. (It might be useful to look at the manual pages of these shells, for example, type: "man csh"). Although most of the commands people type on the prompt are the name of other UNIX programs (such as ls or more), shells recognize some special commands (called internal commands) which are not program names. For example, the exit command terminates the shell, and the cd command changes the current working directory. Shells directly make system calls to execute these commands instead of forking a child process to handle them.

In this assignment, you will develop a **simple shell**. The shell accepts user commands and then executes each command separately. The shell provides the user with a prompt to enter the next Command. One technique for implementing a shell interface is to have the parent process first read what the user enters on the command line and then create a separate child process that performs the Command. Unless otherwise specified, the parent process waits for the child to exit before continuing. However, UNIX shells typically also allow the child process to run in the background - or concurrently - as well by specifying the ampersand (&) at the end of the Command. The separate child process is created using the fork () system call, and the user's Command is executed by using one of the systems calls in the exec () family.

We can summarize the structure of the shell as follows

1. print out a prompt
2. read a line of input from the user
3. parse the line into the program name, and an array of parameters
4. use the fork() system call to spawn a new child process
  - the child process then uses the exec() system call to launch the specified program
  - the parent process (the shell) uses the wait() system call to wait for the child to terminate
5. when the child (i.e. the launched program) finishes, the shell repeats the loop by jumping to 1.

## Task 1: [80 marks] Build a new shell (cshell)

You will develop a command-line interpreter or shell supporting *the environment variables and history of executed commands*. Let us call it **cshell**. The **cshell** will support basic shell functionalities.

**cshell** recognizes the following lines:

- ~~It recognizes lines of the form \$<VAR>=<value>~~
- ~~It recognizes lines of the form <command> <arg0> <arg1> ... <argN>, where <command> is a name of built-in command.~~

**cshell** will also support the following **built-in** commands:

man builtin to read about built-in commands

- ~~exit, the shell terminates on this Command. This must be implemented for a clean exit of the program.~~
- ~~log, the shell prints history of executed commands with time and return code~~
- ~~print, the shell prints argument given to this Command~~
- ~~theme, the shell changes the color of and output~~

When **cshell** takes a **non-built-in** command (like ls, pwd, whoami), it is executed in the child process, and Command's output is printed. **cshell** creates a child process using *fork()* system call, then **cshell** waits for the child process to terminate via *wait()* system call. Child process executes a non-built-in command using *exec()* and its analogues.

You have to create a pipe from the parent process to the child, using *pipe()*. Then you must redirect the standard output (STDOUT\_FILENO) and error output (STDERR\_FILENO) using *dup* or *dup2* to the pipe, and in the parent process, read from the pipe. This is needed to control the output of commands.

So **cshell** should be able to parse the command name and its arguments. Clearly, the hint is that the command name always goes first, and arguments are separated by space.

### Two modes

Our cshell will work in two modes: *interactive mode and script mode*.

The **interactive mode** is activated when **cshell** starts without command line arguments.

Interactive mode is essentially the following loop:

- Print out a prompt
- Read line
- Parse line, and if the line is non-valid, print an error message
- If the line was valid, do what should be done.

### Script mode

For script mode, it should start like `./cshell <filename>`

The script is a file containing a set of lines e.g.

```
.<command1> <arg0> <arg1> ... <argN>
$VAR1=<value1>
```

.....

<commandK> <arg0> <arg1> ... <argN>

.....

In **script mode**, the shell does the following for each line of the file:

- Read the line of the file
- Parse line, if the line is non-valid, print an error message
- If the line was valid, do what should be done.
- The program must exit cleanly after executing in script mode.

You must submit *myscript.txt* showing your input for script mode. The last input must be the log command. For example:

```
$VAR=foo
print $VAR
pwd
whoami
theme red
log
```

**Note :** For the above *myscript.txt*, the program must exit cleanly after execution.

### Environment variables

~~The shell will support the inner environment variables. Each environment variable could be stored in a struct like~~

```
typedef struct {
    char *name;
    char *value;
} EnvVar;
```

~~When **cshell** takes a line of the form \$<VAR>=<value>, it should allocate a new EnvVar struct with name=<VAR> and value=<value>. All environment variables could be stored in some array. EnvVar's name should be unique. If a variable already exists, its value should be updated.~~

~~If some argument takes the form \$<VAR>, **cshell** should look up stored environment variables, find the one with name=<VAR> and substitute argument with environment variable's value. If the corresponding variable does not exist, an error message must be printed.~~

Parsing of lines of the form \$<VAR>=<value> should be simple given that it starts with \$ symbol and variable name and value separated by = sign. Parsing of lines of the form <command> <arg0> <arg1> ... <argN> gets a little more complicated, shell should check if <arg> starts with \$ symbol.

### Built-in commands

We mentioned earlier that the shell would support four built-in commands *exit*, *print*, *theme* and *log*.

1. ~~The *cshell* must start with *cshell\$* (no other characters or welcome messages are allowed)~~
2. ~~The *exit* is simple. It just terminates the shell with a message "Bye!"~~  
~~Note. The output should be exactly same as shown below (e.g. Bye must have exclamation sign, Bye should be case sensitive, no other messages, symbols allowed)~~

```
cshell$ exit
Bye!
```

3. ~~The *print* command takes arguments <arg0> <arg1> ... <argN> and just prints them.~~  
~~Some examples:~~

```
cshell$ print "hazra"
"hazra"
cshell$ print 'hazra'
'hazra'
cshell$ print 123
123
cshell$ print true
true
cshell$ print A[10]
A[10]
```

```
cshell$ print "Good" "Morning"
"Good" "Morning"
```

```
cshell$ $var = "foo"
Variable value expected
cshell$ $var=foo
cshell$ print $var
foo
cshell$ $var="foo"
cshell$ print $var
"foo"
cshell$
```

4. The *log* command history of executed commands with time and return code. So shell should store for each executed command struct like

```
typedef struct {
    char *name;
    struct tm time;
    return value;
} Command;
```

Note: struct tm defined in <time.h>

5. So the *log* command prints an array of such structs.

```

cshell$ log
Thu Jan 27 09:49:33 2022
print 0
Thu Jan 27 09:49:37 2022
print 0
Thu Jan 27 09:49:42 2022
print 0
Thu Jan 27 09:49:45 2022
print 0
Thu Jan 27 09:49:55 2022
print 0
Thu Jan 27 09:50:40 2022
log 0
Thu Jan 27 09:50:46 2022
ls 0
Thu Jan 27 09:50:48 2022
pwd 0

```

6. The *theme* command takes one argument: a name of a colour. And then, the shell using ANSI escape codes changes the colour of its output. Cshell should support three colours (red, green and blue). For any other color, cshell must display a message “**unsupported theme**” (Its case sensitive).

Hint on the steps to get started

- Read the man pages for `fork()`, `execvp()`, `wait ()`, `dup2()`, `open()`, `read()`, `fgets()`, `write()`, `exit()`, `malloc()`, `realloc()`, `strtok()`, `strdup()` system calls
- Write a loop to read a line and split it into tokens
- Write a function to parse tokens
- Come up with how to store environment variables and log information
- Write a function to handle built-in commands and to execute external files (`execvp`)
- Handle reading from file
- Be aware of memory leaks. You can manage time with `<time.h>` library.
- You cannot use any third-party libraries.

**Sample output:**

1. Using the `./cshell` in an interactive mode:

```
cshell$ ls
Makefile
cshell
cshell.c
cshell.h
cshell.o
script.txt
syscall_implementation
test.txt
test_script.txt
cshell$ theme red
cshell$ print "ABC"
"ABC"
cshell$ theme blue
cshell$ print 123
123
cshell$ log
Thu Jan 27 11:04:30 2022
ls 0
Thu Jan 27 11:04:35 2022
theme 0
Thu Jan 27 11:04:42 2022
print 0
Thu Jan 27 11:04:51 2022
theme 0
Thu Jan 27 11:04:55 2022
print 0
cshell$
```

2. Using the ./cshell script.txt in script mode:

```
1 $VAR=foo
2 print $VAR
3 pwd
4 whoami
5 theme red
6 log
7
```

Script.txt

```

himran@TABLET-FCBQI4FM:/mnt/c/Users/hazra/CMPT300demo/A2/cshell_A2/cshell$ ./cshell script.txt
foo
/mnt/c/Users/hazra/CMPT300demo/A2/cshell_A2/cshell
himran
Thu Jan 27 11:08:11 2022
$VAR=foo 0
Thu Jan 27 11:08:11 2022
print 0
Thu Jan 27 11:08:11 2022
pwd 0
Thu Jan 27 11:08:11 2022
whoami 0
Thu Jan 27 11:08:11 2022
theme 0
Bye!
himran@TABLET-FCBQI4FM:/mnt/c/Users/hazra/CMPT300demo/A2/cshell_A2/cshell$ 

```

### Testcases:

Run the ./cshell in an interactive mode:

- Run the built-in commands and verify the output. The commands must include print, theme, log and exit.
- Assign values to any variables in your C program and print the updated values.
- Run non-built-in (pwd, ls) commands and verify the output.

Run the ./cshell in script mode:

- Place all the commands in a .txt file and run cshell bypassing this file
- The result should be the output of all the commands in the file.
- Place an error case in your file. For example, have a line as "failure," and the output should print missing Command.

### Error outputs

```

cshell$ prin "test"
Missing keyword or command, or permission problem
cshell$ prtint2 "test"
Missing keyword or command, or permission problem

```

```

cshell$ theme gree
Unsupported theme: gree
cshell$ theme green
cshell$ 

```

```

himran@TABLET-FCBQI4FM:/mnt/c/Users/hazra/CMPT300demo/A2/cshell_A2/cshell$ ./cshell wrongscript.txt
Unable to read script file: wrongscript.txt

```



## BONUS QUESTION - This is optional.

### **Bonus :** [20 marks] Add a new system call

The task is to add a new system call to the kernel, which will perform the following actions:

- accepts a pointer from the user's application to a string of ASCII characters with codes 32-127 and the length of this string in bytes;
- converts string characters in the range 0x61-0x7A (a-z) to upper case and returns the string.

### Steps on how to add the system call:

1. Make sure you have the source code of the Linux Kernel
2. Create a folder with your C file and a Makefile in the root directory of the kernel sources
3. Add the system call to the system call table
4. Define the macros associated with each system call
  - In arch/x86/include/asm/unistd\_32.h:
    - add the definition for our new system call
    - incremented the value of the macro NR\_SYSCALLS
  - In arch/x86/include/asm/unistd\_64.h add the macro definition
  - In include/linux/syscalls.h add the prototype of the system call
5. The kernel maintains a Makefile as well. Add your directory in the core-y field
6. Compile the kernel

### References:

- <https://arvindraj.wordpress.com/2012/10/05/adding-hello-world-system-call-to-linux/>

- compiling the kernel:

<https://www.howopensource.com/2011/08/how-to-compile-and-install-linux-kernel-3-0-in-ubuntu-11-04-10-10-and-10-04/>

- [Add your own system call to the Linux kernel](#)

For Ubuntu 20.04 users:

<https://dev.to/jasper/adding-a-system-call-to-the-linux-kernel-5-8-1-in-ubuntu-20-04-lts-2ga8>

### Hint on the steps to get started

- You should read how to recompile the kernel and try to do this
- Write system call login in regular function and test it
- Add a new system call

**Now, integrate the system call into the shell:** Add a new command `uppercase <string>` which prints a string modified by your new system call.

### Testcase:

Pass any random string to your program, and the result should print the value in an uppercase on the terminal.

## Submission Instructions:

For A2, the concrete required deliverables are:

- group.txt
- myscript.txt
- cshell.c (with bonus question code if any)
- .h files (if any)
- Makefile
- Bonus question - Also submit a snapshot showing your terminal's output displaying the result of the implemented system call.

You can check the grade file to see your grades before submitting the code. Below is the screenshot.

The screenshot shows a Github Actions interface for a workflow named 'Add files via upload C/C++ CI #3'. The 'build' job is selected and shows a list of steps, all of which are successful. The 'check grade files' step is expanded, showing the following output:

```
1 ▶ Run cat Student_grade.txt
4 Grade for script mode execution : 20
5 Grade for make file execution : 10
6 Grade for ls and who command execution : 5
7 Grade for theme (green + purple) command execution : 10
8 Grade for log command execution : 10
9 Grade for exit + bye command execution : 10
10 Grade for print command execution : 10
11 Total Grade is : 75/80
12 Grade for Bonus Question (Optional) : /20
13 TA NAME :
```

**Note:** On Github actions, **who** command will not return correctly (as it has to check it on local), so dont worry. Make sure that who command work on your machine. If it is working then it will work on our grading machine too.

## Grading Criteria

### Rubric (Out of 80)

Incorrect makefile: -80 (0 grade)

Correct makefile: +10 marks  
Interactive mode - exit +10  
Interactive mode- log + 10  
Interactive mode - print +10  
Interactive mode - theme +10  
Interactive mode - non-built-in command +10  
Script mode - +20  
Bonus (optional)- uppercase +20

## FAQ

1. I am just about to start the assignment 2. Can i still work in csil?

Working in csil is fine for Task 1. But for the bonus question, you will need to setup new syscall on your own virtual environment.